

DDA3020 Machine Learning

Lecture 05 Linear Regression

Jicong Fan
School of Data Science, CUHK-SZ

02/13/2025

Outline

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Notations

- A **set** is an **unordered collection** of unique elements.
- We denote a set as a **calligraphic capital character**, for example, S .
- A set of numbers can be **finite** (include a fixed amount of values). In this case, it is denoted using accolades, for example, $\{1, 3, 18, 23, 235\}$ or $\{x_1, x_2, x_3, \dots, x_d\}$.
- A set can also be **infinite**.

Notations

- A set can be infinite and include all values in some interval.
- If a set includes all values between a and b , including a and b , it is denoted using brackets as $[a, b]$.
- If the set does not include the values a and b , such a set is denoted using parentheses like this: (a, b) .
- For example, the set $[0, 1]$ includes such values as 0, 0.0001, 0.25, 0.784, 0.9995, and 1.0.
- A special set denoted $\mathcal{R}$ (or $R, \mathbb{R}$) includes all real numbers from minus infinity to plus infinity.

Notations

- **Intersection** of two sets:


$$\mathcal{S}_3 \leftarrow \mathcal{S}_1 \cap \mathcal{S}_2$$

$$\text{Example: } \{1, 3, 5, 8\} \cap \{1, 8, 4\} = \{1, 8\}$$

- **Union** of two sets:

$$\mathcal{S}_3 \leftarrow \mathcal{S}_1 \cup \mathcal{S}_2$$


$$\text{Example: } \{1, 3, 5, 8\} \cup \{1, 8, 4\} = \{1, 3, 4, 5, 8\}$$



- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Functions, derivative and gradient

- A **function** is a relation that associates each element x of a **set** $\mathcal{X}$, the **domain** of the function, to a single element y of another **set** $\mathcal{Y}$, the **codomain** of the function.
- A function usually has a name. If the function is called f , this relation is denoted $y = f(x)$ (read f of x), the element x is the **argument or input** of the function, and y is the **value of the function or the output**.
- The symbol that is used for representing the input is the variable of the function (we often say that f is a function of the variable x).



Functions, derivative and gradient

- A **scalar function** can also have vector argument such as, $y = f(\mathbf{x})$, or a **scalar** argument ($y = f(x)$).
- A **vector function**, denoted as $\mathbf{y} = \mathbf{f}(\mathbf{x})$, is a function that returns $\mathbf{y}$, which can have either a vector argument ($\mathbf{y} = \mathbf{f}(\mathbf{x})$) or a **scalar** argument ($\mathbf{y} = \mathbf{f}(x)$).

Functions, derivative and gradient

Notation

- The notation $f : \mathbb{R}^d \rightarrow \mathbb{R}$ means that f is a function that maps real d -vectors to real numbers, *i.e.*, it is a scalar-valued function of d dimension vectors.
- If $\mathbf{x}$ is a d -vector, then $f(\mathbf{x})$, which is a scalar, denotes the value of the function f at $\mathbf{x}$. In the notation $f(\mathbf{x})$, $\mathbf{x}$ is referred to as the argument of the function

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_d)$$

Linear and Affine functions

- To describe a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ we have to specify its value for any possible argument $\mathbf{x} \in \mathbb{R}^d$.
- For example, we can define a function $f : \mathbb{R}^4 \rightarrow \mathbb{R}$ by

$$f(\mathbf{x}) = x_1 + x_2 - 2x_3 - x_4$$

Linear and Affine functions

Another example: The inner product function

- Suppose there is a d -vector. We can define a scalar-valued function f of d -vectors, given by

$$f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} = a_1x_1 + a_2x_2 + \dots + a_dx_d$$

for any d -vector $\mathbf{x}$.

- This function gives the inner product of its d -vector argument $\mathbf{x}$ with some (fixed) d -vector $\mathbf{a}$.
- We can also think of f as forming a weighted sum of the elements of $\mathbf{x}$; the elements of $\mathbf{a}$ give the weights used in forming the weighted sum.

Functions, derivative and gradient

Linear and Affine functions

- A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **linear** if it satisfies the following **two properties**:
 - **Homogeneity**: For any d -vector $\mathbf{x}$ and a scalar α , $f(\alpha\mathbf{x}) = \alpha f(\mathbf{x})$.
 - **Additivity**: For any d -vectors $\mathbf{x}$ and $\mathbf{y}$, $f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$

$$f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} = a_1x_1 + a_2x_2 + \dots + a_dx_d$$

for any d -vector $\mathbf{x}$.

- **Homogeneity** states that **scaling** the (vector) argument is the same as scaling the function value.
- **Additivity** says that **adding** (vector) arguments are the same as adding the function values.

Linear and Affine functions

Superposition and linearity

- The inner product function f defined before satisfies the linearity property

$$\begin{aligned}f(\alpha \mathbf{x} + \beta \mathbf{y}) &= \mathbf{a}^\top (\alpha \mathbf{x} + \beta \mathbf{y}) \\&= \mathbf{a}^\top (\alpha \mathbf{x}) + \mathbf{a}^\top (\beta \mathbf{y}) \\&= \alpha (\mathbf{a}^\top \mathbf{x}) + \beta (\mathbf{a}^\top \mathbf{y}) \\&= \alpha f(\mathbf{x}) + \beta f(\mathbf{y})\end{aligned}$$

for all d -vectors $\mathbf{x}, \mathbf{y}$, and all scalars α, β .

- This property is called **superposition** (which consists of homogeneity and additivity).
- A **function** that satisfies the superposition property is called **linear**

Linear and Affine functions

- If a function f is **linear**, superposition extends to linear **combinations of any number of vectors**:

$$f(\alpha_1 \mathbf{x}_1 + \dots + \alpha_k \mathbf{x}_k) = \alpha_1 f(\mathbf{x}_1) + \dots + \alpha_k f(\mathbf{x}_k)$$

for any d -vectors $\mathbf{x}_1, \dots, \mathbf{x}_k$ and any scalars $\alpha_1, \dots, \alpha_k$

Functions, derivative and gradient

Linear and Affine functions

A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **affine** if and only if it can be expressed as $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} + b$ for some d -vector $\mathbf{a}$ and a scalar b , which is sometimes called the **offset**.

Example:

$$f(\mathbf{x}) = 2.3 - 2x_1 + 1.3x_2 - x_3$$

is affine, with $b = 2.3$, $\mathbf{a} = \begin{bmatrix} -2 \\ 1.3 \\ -1 \end{bmatrix}$.

Functions, derivative and gradient




Figure 2.1 *Left.* The function f is linear. *Right.* The function g is affine, but not linear.

Functions, derivative and gradient

- We say that $f(x)$ has a **local minimum** at $x = c$ if $f(x) \geq f(c)$ for every x in some open interval around $x = c$.
- An **interval** is a set of real numbers with the property that any number that lies between two numbers in the set is also included in the set.
- An **open interval** does not include its endpoints and is denoted using parentheses. For example, $(0, 1)$ means “all numbers greater than 0 and less than 1”
- The minimal value among all the **local minima** is called the **global minimum**. See the illustration in the Figure on the next page.

Functions, derivative and gradient




Figure: Local and global minima of a function. $a < x < b$

Functions, derivative and gradient

max vs arg max

- Given a set of values $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$, the operator $\max_{a \in \mathcal{A}} f(a)$ returns the highest value $f(a)$ for all elements in the set $\mathcal{A}$.
- On the other hand, the operator $\arg \max_{a \in \mathcal{A}} f(a)$ returns the element of the set $\mathcal{A}$ that maximizes $f(a)$.
- Sometimes, when the set is implicit or explicit, we can write

$$\max_a f(a) \quad \text{or} \quad \arg \max_a f(a)$$

- Operator min and arg min operates in a similar manner.
- Note: **arg max** returns a value from the **domain** of the function and **max** returns from the **range (codomain)** of the function

Derivative and Gradient

- A **derivative** f' of a function f is a function or a value that describes how fast f grows (or decreases).
- If the derivative is a **constant** value, like 5 or -3 , then the function grows (or decreases) constantly at any point x of its domain.
- If the derivative f' is **positive at some point x** , then the function f grows at this point.
- If the derivative f' is **negative at some point x** , then the function f decreases at this point.
- The **derivative of zero at x** means that the **function's slope** at x is **horizontal**.

Partial Derivative

- Differentiation of a **scalar** function *w.r.t.* a **vector**
- If $f(\mathbf{w})$ is a **scalar function** of d variables, $\mathbf{w}$ is a $d \times 1$ vector, then differentiation of $f(\mathbf{w})$ *w.r.t.* $\mathbf{w}$ results in a $d \times 1$ vector.

$$\frac{df(\mathbf{w})}{d\mathbf{w}} = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \vdots \\ \frac{\partial f}{\partial w_d} \end{bmatrix}$$

This is referred to as the **gradient** of $f(\mathbf{w})$ and written as $\nabla_{\mathbf{w}} f$.

Functions, derivative and gradient

Partial Derivative

- Differentiation of a **vector** function *w.r.t.* a **vector**
- If $\mathbf{f}(\mathbf{w})$ is a vector function of size $h \times 1$ and $\mathbf{w}$ is a $d \times 1$ vector, then differentiation of $\mathbf{f}(\mathbf{w})$ *w.r.t.* $\mathbf{w}$ results in a $d \times h$ vector.

$$\frac{d\mathbf{f}(\mathbf{w})}{d\mathbf{w}} = \begin{bmatrix} \frac{\partial f_1}{\partial w_1} & \cdots & \frac{\partial f_h}{\partial w_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_1}{\partial w_d} & \cdots & \frac{\partial f_h}{\partial w_d} \end{bmatrix}$$

- This is referred to as the **Jacobian** matrix of $\mathbf{f}(\mathbf{w})$, *i.e.*,

$$\mathbf{J} = \frac{d\mathbf{f}(\mathbf{w})}{d\mathbf{w}},$$
$$\mathbf{J}_{ij} = \frac{\partial f_j}{\partial w_i}.$$

Some Vector-Matrix Differentiation Formulations

$$\frac{d(\mathbf{X}^\top \mathbf{w})}{d\mathbf{w}} = \mathbf{X}, \text{ where } \mathbf{X} \text{ is not a function of } \mathbf{w}$$

$$\frac{d(\mathbf{y}^\top \mathbf{X} \mathbf{w})}{d\mathbf{w}} = \mathbf{X}^\top \mathbf{y}$$

$$\frac{d(\mathbf{w}^\top \mathbf{X} \mathbf{w})}{d\mathbf{w}} = (\mathbf{X} + \mathbf{X}^\top) \mathbf{w}$$

- Note that we adopt the **denominator layout** derivative. If you use the **numerator layout** derivative, then all the above results will be transposed.
- Both types are OK, but keep it **consistent** in all derivatives.
- Please refer to the following wiki page for more details:
https://en.wikipedia.org/wiki/Matrix_calculus

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression**
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Linear regression

Dataset: A collection of m labeled examples $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$, with $\mathbf{x}_i \in \mathcal{X}$ being the d -dimensional feature vector of the i -th example, and $y_i \in \mathcal{Y}$ being a real-valued target.

Linear hypothesis function:

- We want to build a **linear** model, *i.e.*, a linear hypothesis function, as follows

$$f_{w_0, \dots, w_d}(\mathbf{x}) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_d x_d,$$

where w_0 is often called bias and $w_1, \dots, w_d$ are often called coefficients.

- Let $\mathbf{w} = [w_0, w_1, \dots, w_d]^\top$ and rewrite $\mathbf{x} \leftarrow [1, x_1, \dots, x_d]^\top$ (augmented feature vector). Then the linear model is equivalent to

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

- **Note:** $f_{\mathbf{w}}$ is called **linear** due to the linearity *w.r.t.* the parameter vector $\mathbf{w}$, rather than *w.r.t.* the original feature vector $[x_1, \dots, x_d]^\top$.

Task of linear regression:

- Using $f_{\mathbf{w}}$ to approximate the ground-truth target function $t: \mathcal{X} \rightarrow \mathcal{Y}$.
- **Note:** If $\mathcal{Y}$ is a **finite and discrete** set, then the task corresponds to a **classification problem**; if $\mathcal{Y}$ is a **continuous** space, then the task corresponds to a **regression problem**.

Linear regression: deterministic perspective

Learning objective function

- To find the optimal values for $\mathbf{w}$ which **minimizes** the following expression:

$$\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

- In mathematics, the expression we minimize or maximize is called an **objective function**, or, simply, an **objective**.

Linear regression: deterministic perspective

- The expression $(f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$ in the above objective is called the **loss function**. It's a penalty measure of mismatch for example i .
- This particular choice of the loss function is called **squared error loss**.
- All model-based learning algorithms have a loss function and what we do to find the best model is we try to minimize the objective known as the **cost function**.
- In linear regression, the cost function is given by the average loss, also called the **empirical risk**.

Linear regression: probabilistic perspective

- We assume that the relationship between the input variable/feature $\mathbf{x}$ and the output variable y is

$$y = \mathbf{w}^\top \mathbf{x} + e, \text{ where } e \sim \mathcal{N}(0, \sigma^2), \quad (1)$$

where e is called **observation noise** or **residual error**, and it is independent with any specific input $\mathbf{x}$.

- Thus, the output y can also be seen as a random variable, and its conditional probability is formulated as

$$p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \mathbf{x}, \sigma^2) \quad (2)$$

Linear regression: probabilistic perspective

Maximum log-likelihood estimation:

- The parameter $\mathbf{w}$ can be learned by [maximum log-likelihood estimation \(MLE\)](#), given the training dataset $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^m$, as follows

$$\mathbf{w}_{MLE} = \arg \max_{\mathbf{w}} \log \mathcal{L}(\mathbf{w}; D), \quad (3)$$

$$\begin{aligned} \log \mathcal{L}(\mathbf{w}; D) &= \log \left(\prod_{i=1}^m p(y_i | \mathbf{x}_i, \mathbf{w}) \right) = \sum_{i=1}^m \log \mathcal{N}(\mathbf{w}^\top \mathbf{x}_i, \sigma^2) \\ &= -m \log(\sigma(2\pi)^{\frac{1}{2}}) - \frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \mathbf{w}^\top \mathbf{x}_i)^2. \end{aligned} \quad (4)$$

- Removing the constants *w.r.t.* $\mathbf{w}$,

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \frac{1}{2} \sum_{i=1}^m (y_i - \mathbf{w}^\top \mathbf{x}_i)^2, \quad (5)$$

which is [exactly the same with](#) the cost function from the [deterministic](#) perspective.

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Linear regression with analytical solution

Learning (Training)

- Consider the set of feature vector $\mathbf{x}_i$ and target output y_i indexed by $i = 1, \dots, m$, then a linear model $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w} = \sum_{j=0}^d w_j x_j$ can be packed as

$$f_{\mathbf{w}}(\mathbf{X}) \quad \Leftrightarrow \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

Learning model

Learning target vector

$$= \begin{bmatrix} \mathbf{x}_1^\top \mathbf{w} \\ \vdots \\ \mathbf{x}_m^\top \mathbf{w} \end{bmatrix}$$

$$\text{where } \mathbf{x}_i^\top \mathbf{w} = [1, x_1, \dots, x_d]_i \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \text{ and } \mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_m^\top \end{bmatrix}$$

Note: The **bias term** is responsible for **shifting the line/plane** up or down.

Least Squares Regression

- In vector-matrix notation, the squared error function for all the m samples can be written compactly using $\mathbf{e} = \mathbf{X}\mathbf{w} - \mathbf{y}$:

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 \\ &= \sum_{i=1}^m e_i^2 \\ &= \mathbf{e}^\top \mathbf{e} \\ &= (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= (\mathbf{w}^\top \mathbf{X}^\top - \mathbf{y}^\top) (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - \mathbf{w}^\top \mathbf{X}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y} \\ &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y} \end{aligned}$$

Linear Regression

Differentiating $J(\mathbf{w})$ with respect to $\mathbf{w}$ and setting the result to 0 :

$$\begin{aligned}\frac{\partial}{\partial \mathbf{w}} J(\mathbf{w}) &= \mathbf{0} \\ \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y}) &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{X}^\top \mathbf{y} &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X} \mathbf{w} &= 2\mathbf{X}^\top \mathbf{y}\end{aligned}$$

If $\mathbf{X}^\top \mathbf{X}$ is invertible, then

$$\text{Learning: } \hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\text{Prediction: } f_{\mathbf{w}}(\mathbf{X}_{\text{new}}) = \mathbf{X}_{\text{new}} \hat{\mathbf{w}}$$

Linear Regression

Example 1

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^m \quad \begin{aligned} \{x = -9\} &\rightarrow \{y = -6\} \\ \{x = -7\} &\rightarrow \{y = -6\} \\ \{x = -5\} &\rightarrow \{y = -4\} \\ \{x = 1\} &\rightarrow \{y = -1\} \\ \{x = 5\} &\rightarrow \{y = 1\} \\ \{x = 9\} &\rightarrow \{y = 4\} \end{aligned}$$

$$\begin{bmatrix} 1 & -9 \\ 1 & -7 \\ 1 & -5 \\ 1 & 1 \\ 1 & 5 \\ 1 & 9 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix}$$

$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$


(Least squares approximation).

This set of linear equations has NO exact solution.

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible})$$

$$= \begin{bmatrix} 6 & -6 \\ -6 & 262 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ -9 & -7 & -5 & 1 & 5 & 9 \end{bmatrix} \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix} = \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix}$$

Linear regression



$$\hat{y} = \mathbf{X}\hat{\mathbf{w}}$$

$$= \mathbf{X} \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix}$$

$$y = -1.4375 + 0.5625x$$

Prediction:

$$\{x = -1\} \rightarrow \{y = ?\}$$

$$\begin{aligned} \hat{y} &= [1 \ -1] \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix} \\ &= -2 \end{aligned}$$

Linear Regression for one-dimensional examples.

Linear Regression

Example 2

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^m \quad \begin{aligned} \{x_1 = 1, x_2 = 1, x_3 = 1\} &\rightarrow \{y = 1\} \\ \{x_1 = 1, x_2 = -1, x_3 = 1\} &\rightarrow \{y = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 3\} &\rightarrow \{y = 2\} \\ \{x_1 = 1, x_2 = 1, x_3 = 0\} &\rightarrow \{y = -1\} \end{aligned}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix} \quad \mathbf{w}^T \mathbf{X}^T$$

$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$

This set of linear equations has NO exact solution.

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible})$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix} \approx \begin{bmatrix} -0.75 \\ 0.18 \\ 0.93 \end{bmatrix} \text{ (Least squares approximation).}$$

Prediction:

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{y = ?\}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y = ?\}$$

$$\begin{aligned}\hat{y} &= \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.75 \\ 0.18 \\ 0.93 \end{bmatrix} \\ &= \begin{bmatrix} 7.7500 \\ -1.6786 \end{bmatrix}\end{aligned}$$

Linear regression solved by gradient descent

- The linear regression is formulated for the following optimization problem

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w}), \quad J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 = \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \quad (6)$$

- $\mathbf{w}$ can be updated by [gradient descent algorithm](#),

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}, \quad \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \quad (7)$$

where α is called step size or learning rate.

- Does gradient descent always converge to the optimal solution? ([Plot the update trajectory of gradient descent on loss curve](#))

Closed-form solution vs. gradient descent

Closed-form solution $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$	Gradient descent $\mathbf{w} \leftarrow \mathbf{w} - \alpha \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$
No hyper-parameter	Needs to choose α
No need to iterate	Needs many iterations
Complexity $O(d^3 + md^2)$	Complexity $O(T \times md^2)$
Slow if d is very large	Works well when d is large

Thus, you can choose between the above two solutions according to the dimensionality of your training data:

- When the training data is very high-dimensional, *i.e.*, d is very **large**, then it is better to choose **gradient descent algorithm**
- When the training data is not high-dimensional, *i.e.*, d is very **small**, then it is better to choose **closed-form solution**

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Linear regression with single output

When considering the entire set of data indexed by $i = 1, \dots, m$, a linear model $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w} = \sum_{j=0}^d w_j x_j$, where $x_0 \equiv 1$, can be packed as

$$f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w} \leftarrow \text{Scalar function}$$
$$= \begin{bmatrix} \mathbf{x}_1^\top \mathbf{w} \\ \vdots \\ \mathbf{x}_m^\top \mathbf{w} \end{bmatrix} \quad \text{where} \quad \mathbf{x}_i^\top \mathbf{w} = [1, x_{i,1}, \dots, x_{i,d}] \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \quad (8)$$

Primal solution: $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

Note: The **bias term** is responsible for **shifting the line/plane** up or down.

Linear regression with multiple outputs

- When considering the entire set of data indexed by $i = 1, \dots, m$, a linear model $\mathbf{f}_{\mathbf{W}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{W}$ can be packed as

$$\mathbf{f}_{\mathbf{W}}(\mathbf{X}) = \mathbf{X}\mathbf{W}$$

The diagram illustrates the matrix representation of linear regression with multiple outputs. It shows the input matrix $\mathbf{X}$, the weight matrix $\mathbf{W}$, and the output matrix $\mathbf{f}_{\mathbf{W}}(\mathbf{X})$.

Input Matrix $\mathbf{X}$: The first part shows the input matrix $\mathbf{X}$ as a collection of samples. Sample 1 is represented by the row vector $\mathbf{x}_1^T$, and Sample m is represented by the row vector $\mathbf{x}_m^T$. These are stacked vertically to form the matrix $\mathbf{X}$.

Weight Matrix $\mathbf{W}$: The weight matrix $\mathbf{W}$ is shown as a matrix where each row corresponds to a feature. The first row contains the bias term 1 followed by features $x_{1,1}, \dots, x_{1,d}$. The last row contains the bias term 1 followed by features $x_{m,1}, \dots, x_{m,d}$. The matrix is partitioned into three vertical sections: the first section contains the bias terms (1s), the second section contains the weights for the first feature ($w_{0,1}, w_{1,1}, \dots, w_{d,1}$), and the third section contains the weights for the h -th feature ($w_{0,h}, w_{1,h}, \dots, w_{d,h}$).

Output Matrix $\mathbf{f}_{\mathbf{W}}(\mathbf{X})$: The output matrix is shown as a matrix where each row corresponds to a sample's output. Sample 1's output is the row vector $[f_{1,1}, \dots, f_{1,h}]$, and Sample m 's output is the row vector $[f_{m,1}, \dots, f_{m,h}]$. These are stacked vertically to form the matrix $\mathbf{f}_{\mathbf{W}}(\mathbf{X})$. The matrix is partitioned into three vertical sections: the first section contains the outputs for the first feature ($f_{1,1}, \dots, f_{m,1}$), the second section contains the outputs for the h -th feature ($f_{1,h}, \dots, f_{m,h}$), and the third section contains the outputs for the other features. A bracket on the right indicates that the matrix has m rows, and a bracket at the bottom indicates that the matrix has h columns.

Linear regression with multiple outputs

In matrix-matrix notation, the squared error loss function can be written compactly using $\mathbf{E} = \mathbf{X}\mathbf{W} - \mathbf{Y}$:

$$\begin{aligned} J(\mathbf{W}) &= \text{trace}(\mathbf{E}^\top \mathbf{E}) \\ &= \text{trace}[(\mathbf{X}\mathbf{W} - \mathbf{Y})^\top (\mathbf{X}\mathbf{W} - \mathbf{Y})] \end{aligned}$$

If $\mathbf{X}^\top \mathbf{X}$ is invertible, then

$$\begin{aligned} \text{Learning: } \widehat{\mathbf{W}} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} \\ \text{Prediction: } \mathbf{f}_{\mathbf{W}}(\mathbf{X}_{\text{new}}) &= \mathbf{X}_{\text{new}} \widehat{\mathbf{W}} \end{aligned}$$

Assumption: the error terms $\mathbf{e}_k^\top \mathbf{e}_k$ are independent for all $k = 1, \dots, h$

Linear regression with multiple outputs

$$J(\mathbf{W}) = \text{trace}(\mathbf{E}^T \mathbf{E})$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \\ \vdots \\ \mathbf{e}_h^T \end{bmatrix} [\mathbf{e}_1 \quad \mathbf{e}_2 \quad \dots \quad \mathbf{e}_h]\right)$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \mathbf{e}_1 & \mathbf{e}_1^T \mathbf{e}_2 & \dots & \mathbf{e}_1^T \mathbf{e}_h \\ \mathbf{e}_2^T \mathbf{e}_1 & \mathbf{e}_2^T \mathbf{e}_2 & \dots & \mathbf{e}_2^T \mathbf{e}_h \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{e}_h^T \mathbf{e}_1 & \mathbf{e}_h^T \mathbf{e}_2 & \dots & \mathbf{e}_h^T \mathbf{e}_h \end{bmatrix}\right) = \sum_{k=1}^h \mathbf{e}_k^T \mathbf{e}_k$$

Linear regression

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^m \quad \begin{aligned} \{x_1 = 1, x_2 = 1, x_3 = 1\} &\rightarrow \{y_1 = 1, y_2 = 0\} \\ \{x_1 = 1, x_2 = -1, x_3 = 1\} &\rightarrow \{y_1 = 0, y_2 = 1\} \\ \{x_1 = 1, x_2 = 1, x_3 = 3\} &\rightarrow \{y_1 = 2, y_2 = -1\} \\ \{x_1 = 1, x_2 = 1, x_3 = 0\} &\rightarrow \{y_1 = -1, y_2 = 3\} \end{aligned}$$

Example

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} \\ w_{2,1} & w_{2,2} \\ w_{3,1} & w_{3,2} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix}$$

$\mathbf{X}$ $\mathbf{W}$ $\mathbf{Y}$


This set of linear equations has **NO** exact solution.

$$\hat{\mathbf{W}} = \mathbf{X}^\dagger \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible})$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix} = \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix} \quad (\text{Least squares approximation}).$$

Linear regression with multiple outputs

Example:



Prediction:

$$\{x_1 = 1, x_2 = 6, \quad x_3 = 8\} \rightarrow \{y_1 = ?, y_2 = ?\}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y_1 = ?, y_2 = ?\}$$

$$\hat{\mathbf{Y}} = \mathbf{X}_t \hat{\mathbf{W}}$$

$$= \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix}$$

$$= \begin{bmatrix} 7.75 & -7.25 \\ -1.6786 & 3.4643 \end{bmatrix}$$

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Linear regression for **classification**

Dataset:

- We have a collection of m labeled examples $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$, with $\mathbf{x}_i \in \mathcal{X}$ being the d -dimensional feature vector of the i -th example, and $y_i \in \mathcal{Y}$ being a real-valued target.

Linear hypothesis function:

- We want to build a linear model $f_{\mathbf{w}}(\mathbf{x})$, *i.e.*, linear hypothesis function,

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w},$$

where $\mathbf{w} = [w_0, w_1, \dots, w_d]^\top$ is a $(d+1)$ -dimensional vector of parameters, w_0 is the bias parameter, and $\mathbf{x}$ is the $(d+1)$ -dimensional augmented feature vector, *i.e.*, $\mathbf{x} = [1, x_1, \dots, x_d]^\top$.

Task of linear regression:


- Using the linear model $f_{\mathbf{w}}$ to approximate the ground-truth target function $t : \mathcal{X} \rightarrow \mathcal{Y}$.
- Note: If $\mathcal{Y}$ is a **finite and discrete** set, then the task corresponds to a **classification problem**; if $\mathcal{Y}$ is a **continuous** space, then the task corresponds to a regression problem.

Linear regression for classification

Binary Classification: If $\mathbf{X}^\top \mathbf{X}$ is invertible, then

Learning: $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$, $y_i \in \{-1, +1\}, i = 1, \dots, m$

Prediction: $f_{\mathbf{w}}(\mathbf{x}_{new}) = \text{sgn}(\mathbf{x}_{new}^\top \hat{\mathbf{w}})$ (for each row $\mathbf{x}_{new}^\top$ of $\mathbf{X}_{new}$)



Linear regression for classification

Example

$$\{\mathbf{x}_i, y_i\}_{i=1}^m \quad \begin{array}{l} \{x = -9\} \rightarrow \{y = -1\} \\ \{x = -7\} \rightarrow \{y = -1\} \\ \{x = -5\} \rightarrow \{y = -1\} \\ \{x = 1\} \rightarrow \{y = +1\} \\ \{x = 5\} \rightarrow \{y = +1\} \\ \{x = 9\} \rightarrow \{y = +1\} \end{array}$$

$$\begin{bmatrix} 1 & -9 \\ 1 & -7 \\ 1 & -5 \\ 1 & 1 \\ 1 & 5 \\ 1 & 9 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} -1 \\ -1 \\ -1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$


This set of linear equations has **NO** exact solution.

However
($\mathbf{X}^T \mathbf{X}$ is invertible)

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \begin{bmatrix} 6 & -6 \\ -6 & 262 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ -9 & -7 & -5 & 1 & 5 & 9 \end{bmatrix} \begin{bmatrix} -1 \\ -1 \\ -1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix} \quad (\text{Least squares approximation}).$$

Linear regression for classification



$$\hat{y} = \text{sgn}(\mathbf{X}\hat{\mathbf{w}})$$

$$= \text{sgn} \left(\mathbf{X} \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix} \right)$$

$$y = 0.1406 + 0.1406x$$

Prediction:

$$\{x = -2\} \rightarrow \{y = ?\}$$

$$\hat{y} = \text{sgn} \left([1 \ -2] \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix} \right)$$

$$= \text{sgn}(-0.1406) = -1$$

Linear regression for one-dimensional classification.

Linear regression for classification

Linear Methods for Multi-Category Classification:

If $\mathbf{X}^\top \mathbf{X}$ is invertible, then

Learning: $\widehat{\mathbf{W}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}$, $\mathbf{Y} \in \mathbf{R}^{m \times C}$

Prediction: $f_{\mathbf{w}}(\mathbf{x}_{\text{new}}) = \arg \max_{i=1, \dots, C} (\mathbf{x}_{\text{new}}^\top \widehat{\mathbf{W}})$ (for each row $\mathbf{x}_{\text{new}}^\top$ of $\mathbf{X}_{\text{new}}$)

Each row (of $i=1, \dots, m$) in $\mathbf{Y}$ has a one-hot assignment:

e.g., target for class-1 is labelled as $\mathbf{y}_i^\top = [1, 0, 0, \dots, 0]$ for the i th sample,

target for class-2 is labelled as $\mathbf{y}_i^\top = [0, 1, 0, \dots, 0]$ for the i th sample,

target for class-C is labelled as $\mathbf{y}_i^\top = [0, 0, \dots, 0, 1]$ for the i th sample.

Linear regression for classification

Example

$$\{\mathbf{x}_i, y_i\}_{i=1}^m \quad \begin{aligned} \{x_1 = 1, x_2 = 1, x_3 = 1\} &\rightarrow \{y_1 = 1, y_2 = 0, y_3 = 0\} \\ \{x_1 = 1, x_2 = -1, x_3 = 1\} &\rightarrow \{y_1 = 0, y_2 = 1, y_3 = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 3\} &\rightarrow \{y_1 = 1, y_2 = 0, y_3 = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 0\} &\rightarrow \{y_1 = 0, y_2 = 0, y_3 = 1\} \end{aligned}$$

$$\begin{array}{c} \begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \\ \mathbf{X} \end{array} \begin{array}{c} \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \end{bmatrix} \\ \mathbf{W} \end{array} = \begin{array}{c} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \\ \mathbf{Y} \end{array}$$

This set of linear equations has **NO** exact solution.

$$\widehat{\mathbf{W}} = \mathbf{X}^{\dagger} \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible}) \quad (\text{Least squares approximation})$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 0.5 & 0.5 \\ 0.2857 & -0.5 & 0.2143 \\ 0.2857 & 0 & -0.2857 \end{bmatrix}$$

Linear regression for classification

Example

Prediction:

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{ \text{class 1, 2, or 3 ?} \}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{ \text{class 1, 2, or 3 ?} \}$$

$$\hat{\mathbf{Y}} = \mathbf{X}_t \hat{\mathbf{W}} = \arg \max_{i=1,\dots,C} \left(\begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} 0 & 0.5 & 0.5 \\ 0.2857 & -0.5 & 0.2143 \\ 0.2857 & 0 & -0.2857 \end{bmatrix} \right)$$

$$= \arg \max_{i=1,\dots,C} \left(\begin{bmatrix} 4 & -2.50 & -0.50 \\ -0.2587 & 0.50 & 0.7857 \end{bmatrix} \right)$$

Position of the largest
number determines
the class label

$$= \begin{bmatrix} 1 \\ 3 \end{bmatrix} \begin{matrix} \rightarrow \text{Class-1} \\ \rightarrow \text{Class-3} \end{matrix}$$

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
 - Deterministic perspective
 - Probabilistic perspective
- 4 Learning of linear regression
 - Analytical solution
 - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
 - Ridge regression
 - Polynomial regression
 - Lasso regression
 - Robust linear regression

Ridge regression

Motivation 1:

Recall the learning computation: $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$.

We cannot guarantee that the matrix $\mathbf{X}^\top \mathbf{X}$ is invertible.

Ridge regression

Here on, we shall focus on single output y in derivations in the sequel.

$$\min_{\mathbf{w}} \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}}, \text{ where } \tilde{\mathbf{w}} = \hat{\mathbf{I}}_d \mathbf{w} = [0, w_1, w_2, \dots, w_d]^\top,$$

$\hat{\mathbf{I}}_d \in \mathbb{R}^{(d+1) \times (d+1)}$ is defined by setting the $(1, 1)$ entry in the $d + 1$ dimensional identity matrix $\hat{\mathbf{I}}_{d+1}$ as 0. **Note:** The bias/offset w_0 is **NOT** included in the ℓ_2 regularization term, as it just affects the function's height, not its complexity.

Linear Model: $\min_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}}$


$$\begin{aligned} \frac{\partial}{\partial \mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}} &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{X}^\top \mathbf{y} + 2\lambda \hat{\mathbf{I}}_d \mathbf{w} &= \mathbf{0} \\ \Rightarrow \mathbf{X}^\top \mathbf{X}\mathbf{w} + \lambda \hat{\mathbf{I}}_d \mathbf{w} &= \mathbf{X}^\top \mathbf{y} \\ \Rightarrow (\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d) \mathbf{w} &= \mathbf{X}^\top \mathbf{y} \\ \Rightarrow \mathbf{w} &= (\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d)^{-1} \mathbf{X}^\top \mathbf{y} \end{aligned}$$

Note that $(\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d)$ is **guaranteed to be invertible**, given $\lambda > 0$.


Ridge regression

Motivation 2:

- **Overfitting** is an important challenge for linear regression, as shown below. Note: M in the figure denotes the degree of polynomial hypothesis function.
- If overfitting, the prediction performance on testing data will be very poor. How to alleviate overfitting?



Ridge regression



- Let's see one simple example, we use a **polynomial function** (introduced later) with 14 degrees to fit $m = 21$ data points. The learned curve is very “wiggly” (see above).
- The parameter values of this curve are as follows
 $6.56, -36.934, -109.25, 543.452, 1022.561, -3046.224, -3768.013, 8524.54, 6607.897, -12640.058, -5530.188, 9479.73, 1774, 639, -2821.526$
- There are many large positive/negative values, such that a small change of features could lead to a significant change of output.

Ridge regression

- How to get smaller parameter values?
- We can assume that the parameter $\mathbf{w}$ (excluding the bias b) follow a **zero-mean Gaussian prior**

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \tau^2\mathbf{I}) \quad (9)$$

- For clarity, **we omit the bias w_0 in the following derivation.**
- Utilizing this prior, we obtain the **maximum a posteriori (MAP)** estimation

$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} \left[\sum_{i=1}^m \log p(y_i|\mathbf{x}_i, \mathbf{w}) + \log p(\mathbf{w}) \right] \quad (10)$$

$$= \arg \max_{\mathbf{w}} \left[\sum_{i=1}^m \log \mathcal{N}(\mathbf{x}_i^\top \mathbf{w}, \sigma^2) + \log \mathcal{N}(\mathbf{w}|\mathbf{0}, \tau^2\mathbf{I}) \right] \quad (11)$$


$$\equiv \arg \min_{\mathbf{w}} \left[\sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 \right]. \quad (12)$$

- The corresponding closed-form solution is given by

$$\mathbf{w}_{MAP} = (\lambda\mathbf{I} + \mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}. \quad (13)$$

Ridge regression

- The above method is also known as **ridge regression**, or **penalized least squares**.
- In general, adding a Gaussian prior to the parameters of a model to encourage them to be small is called ℓ_2 **regularization** or **weight decay**.
- As shown below, when we set a larger λ , *i.e.*, more weight on the prior, the resulting curve will be smoother.





Polynomial regression

Motivation

- Some data may **be not linearly separated**, such as the classic **XOR** data, as shown on the bottom left.
- Consequently, the linear regression model doesn't work.
- To tackle it, we could project the original data to the **monomial** axis x_1x_2 .
- Then, the XOR becomes **linearly separated**, as shown on the bottom right.
- Accordingly, we can design a novel linear regression model, as follows

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + w_{12}x_1x_2 + w_{11}x_1^2 + w_{22}x_2^2$$



Polynomial regression

Polynomial expansion

- The linear model $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w}$ can be written as

$$f_{\mathbf{w}}(\mathbf{x}) = \sum_{i=0}^d x_i w_i = w_0 + \sum_{i=1}^d x_i w_i.$$

- By including terms involving the products of pairs of components of $\mathbf{x}$, we obtain a **quadratic** model:

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j.$$

Polynomial regression

- In general:

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j + \sum_{i=1}^d \sum_{j=1}^d \sum_{k=1}^d w_{ijk} x_i x_j x_k + \dots$$

Remarks

- For high dimensional d and high *polynomial order*, the number of polynomial terms becomes explosive! (In fact, this number of terms grows exponentially.)
- Hence, for high dimensional problems, polynomials of order larger than 3 are seldom used.

Polynomial regression

Linear model with basis expansion $\phi(\mathbf{x})$

$$\begin{aligned} f_{\mathbf{w}}(\mathbf{x}) &= w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j + \sum_{i=1}^d \sum_{j=1}^d \sum_{k=1}^d w_{ijk} x_i x_j x_k + \dots \\ &= \phi(\mathbf{x})^\top \mathbf{w}, \end{aligned}$$

where

$$\begin{aligned} \phi(\mathbf{x}) &= [1, x_1, \dots, x_d, \dots, x_i x_j, \dots, x_i x_j x_k, \dots]^\top, \\ \mathbf{w} &= [w_0, w_1, \dots, w_d, \dots, w_{ij}, \dots, w_{ijk}, \dots]^\top. \end{aligned}$$

Note: $f_{\mathbf{w}}(\mathbf{x})$ is still a linear function *w.r.t.* $\mathbf{w}$, rather than $\mathbf{x}$. Thus, it is still a **linear model**.

Extending to the case of m data points, *i.e.*, $\mathbf{X} = [\mathbf{x}_1^\top; \dots; \mathbf{x}_m^\top] \in \mathbb{R}^{m \times (d+1)}$, the basis expansion is presented by

$$\mathbf{P}(\mathbf{X}) = [\phi(\mathbf{x}_1)^\top; \dots; \phi(\mathbf{x}_m)^\top] \in \mathbb{R}^{m \times |\mathbf{w}|}.$$

Polynomial regression

Example

2nd order polynomial model

$$\begin{aligned}\{\mathbf{x}_i, y_i\}_{i=1}^m & \quad \{x_1 = 0, x_2 = 0\} \rightarrow \{y = 0\} \\ & \quad \{x_1 = 1, x_2 = 1\} \rightarrow \{y = 1\} \\ & \quad \{x_1 = 1, x_2 = 0\} \rightarrow \{y = 2\} \\ & \quad \{x_1 = 0, x_2 = 1\} \rightarrow \{y = 3\}\end{aligned}$$

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + w_{12}x_1x_2 + w_{11}x_1^2 + w_{22}x_2^2$$

$$\begin{aligned}&= \begin{bmatrix} 1 & x_1 & x_2 & x_1x_2 & x_1^2 & x_2^2 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_{12} \\ w_{11} \\ w_{22} \end{bmatrix} \\ \mathbf{P} &= \begin{bmatrix} 1 & x_1 & x_2 & x_1x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1x_2 & x_1^2 & x_2^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 \end{bmatrix}\end{aligned}$$

Polynomial regression

Ridge regression with **original features $\mathbf{X}$** :

$$\text{Learning: } \hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$$

$$\text{Prediction: } f_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{w}}$$

Ridge regression with **basis expansion $\mathbf{P}(\mathbf{X})$** :

$$\text{Learning: } \hat{\mathbf{w}} = (\mathbf{P}^\top \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^\top \mathbf{y}$$

$$\text{Prediction: } f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \mathbf{P}_{new} \hat{\mathbf{w}}$$

Polynomial regression

For Regression Applications

- Learning: $\hat{\mathbf{w}} = (\mathbf{P}^\top \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^\top \mathbf{y}$, where $\mathbf{y}$ is continuous
- Prediction: $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \mathbf{P}_{new} \hat{\mathbf{w}}$

For Classification Applications

- Learn **discrete** valued $\mathbf{y}$ (binary) or $\mathbf{Y}$ (one-hot)
- Binary Prediction: $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \text{sgn}(\mathbf{P}_{new} \hat{\mathbf{w}})$ if $y \in \{-1, +1\}$
- Multi-Category Prediction: $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \text{argmax}_{i=1, \dots, C} (\mathbf{P}_{new} \hat{\mathbf{w}})$

Polynomial regression

Example (Cont'd)

$$\{\mathbf{x}_i, y_i\}_{i=1}^m \quad \begin{aligned} \{x_1 = 0, x_2 = 0\} &\rightarrow \{y = -1\} \\ \{x_1 = 1, x_2 = 1\} &\rightarrow \{y = -1\} \\ \{x_1 = 1, x_2 = 0\} &\rightarrow \{y = +1\} \\ \{x_1 = 0, x_2 = 1\} &\rightarrow \{y = +1\} \end{aligned}$$

2nd order polynomial model

$$\mathbf{P} = \begin{bmatrix} 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 \end{bmatrix}$$

$$\hat{\mathbf{w}} = \mathbf{P}^\top (\mathbf{P}\mathbf{P}^\top)^{-1} \mathbf{y} \quad (\text{note: under determined linear system})$$

$$= \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 6 & 3 & 3 \\ 1 & 3 & 3 & 1 \\ 1 & 3 & 1 & 3 \end{bmatrix}^{-1} \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} = \begin{bmatrix} -1 \\ 1 \\ 1 \\ -4 \\ 1 \\ 1 \end{bmatrix}$$

Polynomial regression

Example (Cont'd)

Prediction:

Test point 1: $\{x_1 = 0.1, x_2 = 0.1\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

Test point 2: $\{x_1 = 0.9, x_2 = 0.9\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

Test point 3: $\{x_1 = 0.1, x_2 = 0.9\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$


Test point 4: $\{x_1 = 0.9, x_2 = 0.1\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

$$\hat{y} = \mathbf{P}_t \hat{\mathbf{w}}$$

$$\hat{y} = \text{sgn} \left(\begin{bmatrix} 1 & 0.1 & 0.1 & 0.01 & 0.01 & 0.01 \\ 1 & 0.9 & 0.9 & 0.81 & 0.81 & 0.81 \\ 1 & 0.1 & 0.9 & 0.09 & 0.01 & 0.81 \\ 1 & 0.9 & 0.1 & 0.09 & 0.81 & 0.01 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \\ 1 \\ -4 \\ 1 \\ 1 \end{bmatrix} \right)$$

$$= \text{sgn} \left(\begin{bmatrix} -0.82 \\ -0.82 \\ 0.46 \\ 0.46 \end{bmatrix} \right)$$

$$= \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix} \begin{array}{l} \dashrightarrow \text{Class } -1 \\ \dashrightarrow \text{Class } -1 \\ \dashrightarrow \text{Class } +1 \\ \dashrightarrow \text{Class } +1 \end{array}$$



Lasso regression

- We can replace the Gaussian prior by a **Laplacian prior**, *i.e.*,

$$p(\mathbf{w}) = \text{Lap}(\mathbf{w}|\mathbf{0}, \lambda) = \frac{1}{2\lambda} \exp\left(-\frac{\|\mathbf{w}\|_1}{\lambda}\right), \quad (14)$$

- The combination of the Gaussian distribution of $p(y|\mathbf{x}, \mathbf{w})$ and the **Laplacian prior**, leading to


$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} \left[\sum_{i=1}^m \log p(y_i|\mathbf{x}_i, \mathbf{w}) + \log p(\mathbf{w}) \right] \quad (15)$$

$$= \arg \max_{\mathbf{w}} \left[\sum_{i=1}^m \log \mathcal{N}(\mathbf{w}^\top \mathbf{x}_i, \sigma^2) + \text{Lap}(\mathbf{w}|\mathbf{0}, b) \right] \quad (16)$$

$$\equiv \arg \min_{\mathbf{w}} \left[\sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 + \alpha \|\mathbf{w}\|_1 \right]. \quad (17)$$


Lasso regression

- It is **Lasso regression**, and the regularization is called ℓ_1 **regularization**. It will encourage the sparse parameters.
- As shown below, when we set a larger α , *i.e.*, more weight on the prior, the resulting curve will be smoother.



Geometry of Ridge and Lasso regression

- Geometry of Ridge and Lasso regression. Which one is Ridge?





Robust linear regression

- When there is a few outliers in the training data D , which are far from most other points, then learned parameters $\mathbf{w}_{MLE}$ will be significantly influenced, leading to a very poor fitting.
- Let's see the loss curve of the residual sum of squares (RSS),

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2. \quad (18)$$

- The error increases quadratically along with the residual. To minimize such a large error, the linear model will be significantly changed.
- How to alleviate the significant influence of outliers?




Robust linear regression

- We adopt the ℓ_1 loss to replace the ℓ_2 loss, as follows

$$J(\mathbf{w}) = \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i|. \quad (19)$$

- The curves of ℓ_1 and ℓ_2 losses are shown as follows.
- When the residual is large, the ℓ_1 loss is much smaller than the ℓ_2 loss, such that the influence of outliers could be alleviated.



Robust linear regression



- Actually, the above ℓ_1 loss can also be derived from the probabilistic perspective, by assuming that

$$p(y|\mathbf{x}, \mathbf{w}, b) = \text{Lap}(y|\mathbf{w}^\top \mathbf{x}, b) \propto \exp\left(-\frac{1}{b}|y - \mathbf{w}^\top \mathbf{x}|\right) \quad (20)$$

- Applying the maximum log-likelihood estimation (MLE), we will obtain

$$\mathbf{w}_{MLE} = \arg \max_{\mathbf{w}} \log \mathcal{L}(\mathbf{w}; D) = \arg \max_{\mathbf{w}} \sum_i^m \log p(y_i|\mathbf{x}_i, \mathbf{w}) \quad (21)$$

$$\equiv \arg \min_{\mathbf{w}} \frac{1}{b} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (22)$$



Robust linear regression

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (23)$$

- However, the ℓ_1 loss function is **non-differentiable**. The gradient descent algorithm cannot be adopted.
- We can transform it to a **linear program**, as follows

$$\begin{aligned} \min_{\mathbf{w}, \mathbf{t}} \quad & \sum_{i=1}^m t_i \\ \text{s.t.} \quad & -t_i \leq \mathbf{x}_i^\top \mathbf{w} - y_i \leq t_i, 1 \leq i \leq m. \end{aligned} \quad (24)$$

Please refer to:

<https://math.stackexchange.com/questions/1639716/how-can-l-1-norm-minimization-with>

Robust linear regression

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (25)$$

- We can also utilize the following equation:

$$|a| = \min_{\mu > 0} \frac{1}{2} \left(\frac{a^2}{\mu} + \mu \right) \quad (26)$$

- Then, the above ℓ_1 minimization problem (25) can be reformulated as follows

$$\min_{\mathbf{w}} \min_{\mu_1, \dots, \mu_m > 0} \frac{1}{2} \left(\frac{(\mathbf{x}_i^\top \mathbf{w} - y_i)^2}{\mu_i} + \mu_i \right). \quad (27)$$

- It can be iteratively and alternatively optimized as follows:
 - Given $\mathbf{w}$, $\mu_i = |\mathbf{x}_i^\top \mathbf{w} - y_i|, i = 1, \dots, m$
 - Given μ , $\mathbf{w} = \arg \min_{\mathbf{w}} \sum_{i=1}^m \frac{1}{2} (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 / \mu_i$
- It is called **iteratively reweighted least squares** method.

Summary of different variants of linear regressions

Note that the uniform distribution will not change the mode of the likelihood.

Thus, MAP estimation with a uniform prior corresponds to MLE.

$p(y \mathbf{x}, \mathbf{w})$	$p(\mathbf{w})$	regression method
Gaussian	Uniform	Least squares
Gaussian	Gaussian	Ridge regression
Gaussian	Laplace	Lasso regression
Laplace	Uniform	Robust regression
Student	Uniform	Robust regression

